

Malik Hassan Raza

+49-152-10737287 • raza.malikhassan55@gmail.com

GenAI Engineer with **4+ years of experience** building **production-grade LLM applications, RAG systems, and AI data pipelines**. Specialized in **LLM orchestration, retrieval architectures, prompt engineering, and evaluation workflows** that improve reasoning quality, reliability, and user experience. Strong background in **Python, LangChain/LangGraph, vector databases, cloud infrastructure, and MLOps**, with a proven ability to translate business needs into scalable GenAI solutions.

WORK EXPERIENCE

Mistral AI

03/2026 – Present

AI Engineer (LLM Evaluations)

- Designed and implemented challenging coding benchmarks to evaluate LLMs across multiple programming languages including Python, TypeScript, Rust, C, C++, and C#.
- Reviewed and analyzed model-generated code, open-source pull requests, and code reviews for correctness, performance, readability, and adherence to project specifications.
- Identified and documented model failures, edge cases, and reasoning gaps, contributing to structured improvement of AI outputs.
- Conducted head-to-head comparisons between private and external LLMs, providing detailed technical feedback to guide model fine-tuning and RLHF pipelines.
- Built and maintained coding environments to support multi-language evaluation tasks and reinforcement learning workflows.
- Applied structured evaluation criteria consistently to support benchmarking, model alignment, and quality control.
- Mentored team members and shared best practices for code evaluation and annotation, improving overall workflow efficiency.

Turing

01/2023 – 10/2025

GenAI Engineer (LLMs, Python)

Remote

- Designed and deployed **LLM-powered systems** for training, evaluation, and internal knowledge access, supporting large-scale GenAI initiatives.
- Led a **pod of 9 GenAI engineers** within a 600+ person organization, improving delivery velocity, data quality, and cross-team alignment.
- Scaled production to **70–80 high-quality LLM trajectories per week**, accelerating model iteration across code generation, reasoning, and conversational tasks.
- Built an **automated LLM-based data quality analysis pipeline** using Python, REST APIs, and GPT-4 to detect systemic errors and quality gaps missed by manual review.
- Engineered an internal **RAG-powered chatbot** using Python, LangChain, and **Qdrant**, reducing Slack message volume by **30%** and enabling instant access to SOPs and documentation.
- Defined **prompting guidelines, evaluation rubrics, and validation workflows** that ensured consistent outputs and **100% client requirement compliance**.
- Partnered with ML researchers and product teams to align **evaluation metrics with benchmarking goals**, directly contributing to measurable model improvements.
- Supported **SFT and RLHF pipelines** by producing high-quality, task-specific datasets that strengthened model alignment and reasoning quality.
- Partnered with cross-functional teams to deliver measurable gains in model performance and product value.

Systems Limited

07/2021 – 01/2023

Junior Consultant Data Analytics

Islamabad

- Built high-performance ETL pipelines using IBM DataStage, improving data ingestion efficiency by 30% while preserving end-to-end data quality for a major telecom client.

- Developed Power BI dashboards and optimized SQL queries, reducing reporting time by 40% and enabling faster data-driven decisions across business units.
- Implemented automated validation workflows with Python/Alteryx, achieving 98%+ post-migration data accuracy, strengthening internal controls and trust in migrated data.

SKILLS

Programming: Alteryx, Git/Github, LangChain, LangGraph, MySQL, Numpy, Pandas, Pinecone, Python, PyTorch, Qdrant, REST API, SQL (analytics, optimization)

ML & LLM Data Systems: GenAI (Prompt Engineering, Custom GPTs, OpenAI APIs), LLMs, RAG (Retrieval Augmented Generation), Vector Databases, Vibe Coding (Cursor, Copilot, Antigravity)

EDUCATION

Bachelor of Science in Computer Science (BSCS)

National University of Sciences and Technology

Islamabad (Pakistan) • 06/2021

PROJECTS

BookFlow

- Building a **GenAI SaaS platform** that transforms creators' existing content into professionally structured books while preserving individual writing voice.
- Designed **RAG pipelines** for style learning, semantic retrieval, outline generation, and chapter-by-chapter manuscript creation.
- Implemented embedding, retrieval, and orchestration workflows using **LangChain, Pinecone, PostgreSQL, and AWS**.
- Integrated LLM APIs for long-form generation, revision, and consistency control.

Tech Stack: React, TypeScript, Node.js, PostgreSQL, Vector DB (Pinecone), Anthropic Claude API, LangChain, Tailwind CSS, AWS, Supabase, Stripe.

Membership Inference Attacks using Shadow Models

- Engineered and scaled shadow models to extract membership information from target classifiers, achieving $\text{TPR@FPR}=0.144$ and $\text{AUC}=0.655$.
- Developed feature-rich attack models using full softmax vectors, entropy, margin, and loss metrics for improved inference performance.
- Implemented dynamic data splitting and scaling strategies with up to 30 shadow models to simulate varied training conditions.
- Leveraged modular pipelines from the yonsei-sslab/MIA framework for streamlined experimentation and reproducibility.

CERTIFICATIONS

Data Engineering Bootcamp

abcData

What is Data Science?

Coursera

Tools for Data Science

Coursera

Python for Data Science, AI & Development

Coursera

Alteryx Micro Credential Foundation

Alteryx